

CS 5350/6350: Machine Learning Fall 2018

Homework 2

Handed out: 13 September, 2018

Due date: 27 September, 2018

General Instructions

- You are welcome to talk to other members of the class about the homework. I am more concerned that you understand the underlying concepts. However, you should write down your own solution. Please keep the class collaboration policy in mind.
- Feel free discuss the homework with the instructor or the TAs.
- Your written solutions should be brief and clear. You need to show your work, not just the final answer, but you do *not* need to write it in gory detail. Your assignment should be **no more than 10 pages**. Every extra page will cost a point.
- Handwritten solutions will not be accepted.
- The homework is due by midnight of the due date. Please submit the homework on Canvas.
- Some questions are marked **For 6350 students**. Students who are registered for CS 6350 should do these questions. Students who have registered for CS 5350 are welcome to do this question for extra credit.

1 Warm up: Linear Classifiers and Boolean Functions

[10 points] Please indicate if each of the following boolean functions is linearly separable. If it is linearly separable, write down a linear threshold unit equivalent to it.

1. [2 points] $\neg x_1 \vee x_2 \vee \neg x_3$
2. [2 points] $(x_1 \vee x_2) \wedge x_3$
3. [2 points] $(x_1 \wedge \neg x_2) \vee \neg x_3$
4. [2 points] $x_1 \text{ xor } x_2 \text{ xor } x_3$
5. [2 points] $x_1 \wedge \neg x_2 \wedge x_3$

2 Feature transformations

[10 points] Consider the concept class C consisting of functions f_r defined by a radius r as follows:

$$f_r(x_1, x_2) = \begin{cases} +1 & 17x_1^4 - 16x_2^3 \leq r \\ -1 & \text{otherwise} \end{cases}$$

Note that the hypothesis class is *not* linearly separable in \mathbb{R}^2 .

Construct a function $\phi(x_1, x_2)$ that maps examples to a new space, such that the positive and negative examples are linearly separable in that space. The answer to this question should consist of two parts:

1. A function ϕ that maps examples to a new space.
2. A proof that in the new space, the positive and negative points are linearly separated. You can show this by producing such a hyperplane in the new space (i.e. find a weight vector \mathbf{w} and a bias b such that $\mathbf{w}^T \phi(x_1, x_2) \geq b$ if, and only if, $f_r(x_1, x_2) = +1$).

Hint: The feature transformation ϕ should not depend on r .

3 Mistake Bound Model of Learning

In both the questions below, we will consider functions defined over n Boolean features. That is, each example in our learning problem is a n -dimensional vector from $\{0, 1\}^n$. We will use the symbol \mathbf{x} to denote an example and \mathbf{x}_i denotes its i^{th} element. (We will assume that there is no noise involved.)

For all questions below, it is not enough to just state the answer. You need to justify your answer with a short proof.

1. Consider the concept class \mathcal{C}_1 defined as follows: Each element of \mathcal{C}_1 is defined using a fixed instance $\mathbf{z} \in \{0, 1\}^n$ as follows:

$$f_{\mathbf{z}}(\mathbf{x}) = \begin{cases} 1 & \mathbf{x} = \mathbf{z} \\ 0 & \mathbf{x} \neq \mathbf{z}. \end{cases}$$

That is, the function $f_{\mathbf{z}}$ predicts 1 if, and only if, the input to the function is \mathbf{z} .

Our goal is to come up with a mistake bound algorithm that will learn any function $f \in \mathcal{C}_1$.

- (a) [5 points] Determine $|\mathcal{C}_1|$, the size of concept class.
- (b) [15 points] Write a mistake bound learning algorithm for this concept class that makes no more than *one* mistake on any sequence of examples presented to it. Please write the algorithm concisely in the form of pseudocode. Prove the mistake bound for this algorithm.

2. Suppose we have a concept class \mathcal{C}_2 that consists of exactly n functions $\{f_1, f_2, \dots, f_n\}$, where each function f_i is defined as follows:

$$f_i(\mathbf{x}) = \mathbf{x}_i.$$

That is, the function f_i returns the value of the i^{th} feature.

- (a) [5 points] How many mistakes will the algorithm **CON** from class make on any function from this concept class?
- (b) [5 points] How many mistakes will the Halving algorithm make on any function from this concept class?

4 The Perceptron Algorithm and its Variants

For this question, you will experiment with the Perceptron algorithm and some variants on a data set.

4.1 The task and data

We will be using the Diabetic Retinopathy dataset from the UCI Machine Learning repository ¹. The dataset consists of features extracted from images and the goal is to predict whether an image contains signs of diabetic retinopathy or not.

Using this labeled data, we want to build a classifier that identifies whether a new retinal image shows signs of diabetic retinopathy or not.

The data has been preprocessed into a standard format. Use the training/development/test files called `diabetes.train`, `diabetes.dev` and `diabetes.test`. These files are in the LIB-SVM format, where each row is a single training example. The format of the each row in the data is:

`<label> <index1>:<value1> <index2>:<value2> ...`

Here `<label>` denotes the label for that example. The rest of the elements of the row is a sparse vector denoting the feature vector. For example, if the original feature vector is $[0, 0, 1, 2, 0, 3]$, this would be represented as `3:1 4:2 6:3`. That is, only the non-zero entries of the feature vector are stored.

4.2 Algorithms

You will implement several variants of the Perceptron algorithm. Note that each variant has different hyper-parameters, as described below. Use 5-fold cross-validation to identify the best hyper-parameters as you did in the previous homework. To help with this, we have split the training set into five parts `training00.data`–`training04.data` in the folder `CVSplits`.

¹<https://archive.ics.uci.edu/ml/datasets/Diabetic+Retinopathy+Debrecen+Data+Set>

1. **Simple Perceptron:** Implement the simple batch version of Perceptron as described in the class. Use a fixed learning rate η chosen from $\{1, 0.1, 0.01\}$. An update will be performed on an example (\mathbf{x}, y) if $y(\mathbf{w}^T \mathbf{x} + b) < 0$ as:

$$\mathbf{w} \leftarrow \mathbf{w} + \eta y \mathbf{x},$$

$$b \leftarrow b + \eta y.$$

Hyper-parameter: Learning rate $\eta \in \{1, 0.1, 0.01\}$

Two things bear additional explanation.

- (a) First, note that in the formulation above, the bias term b is explicitly mentioned. This is because the features in the data do not include a bias feature. Of course, you could choose to add an additional constant feature to each example and not have the explicit extra b during learning. (See the class lectures for more information.) However, here, we will see the version of Perceptron that explicitly has the bias term.
- (b) Second, in this specific case, if \mathbf{w} and b are initialized with zero, then the fixed learning rate will have no effect. To see this, recall the Perceptron update from above.

Now, if \mathbf{w} and b are initialized with zeroes and a fixed learning rate η is used, then we can show that the final parameters will be equivalent to having a learning rate 1. The final weight vector and the bias term will be scaled by η compared to the unit learning rate case, which does not affect the sign of $\mathbf{w}^T \mathbf{x} + b$.

To avoid this, you should initialize the all elements of the weight vector \mathbf{w} and the bias to a small random number between -0.01 and 0.01.

2. **Decaying the learning rate:** Instead of fixing the learning rate, implement a version of the Perceptron algorithm whose learning rate decreases as $\frac{\eta_0}{1+t}$, where η_0 is the starting learning rate, and t is the time step. Note that t should keep increasing across epochs. (That is, you should initialize t to 0 at the start and keep incrementing it each time a new example is encountered.)

Hyper-parameter: Initial learning rate $\eta_0 \in \{1, 0.1, 0.01\}$

3. **Margin Perceptron:** This variant of Perceptron will perform an update on an example (\mathbf{x}, y) if $y(\mathbf{w}^T \mathbf{x} + b) < \mu$, where μ is an additional positive hyper-parameter, specified by the user. Note that because μ is positive, this algorithm can update the weight vector even when the current weight vector does not make a mistake on the current example. You need to use the decreasing learning rate as before.

Hyper-parameters:

- (a) Initial learning rate $\eta_0 \in \{1, 0.1, 0.01\}$
- (b) Margin $\mu \in \{1, 0.1, 0.01\}$

Note: When there is more than one hyper-parameter to cross-validate, you need to consider all combinations of the hyper-parameters. In this case, you will need to perform cross-validation for all pairs (η_0, μ) from the above sets.

4. **Averaged Perceptron** Implement the averaged version of the original Perceptron algorithm from the first question. Recall from class that the averaged variant of the Perceptron asks you to keep two weight vectors (and two bias terms). In addition to the original parameters (\mathbf{w}, b) , you will need to update the averaged weight vector \mathbf{a} and the averaged bias b_a as:

$$(a) \quad \mathbf{a} \leftarrow \mathbf{a} + \mathbf{w}$$

$$(b) \quad b_a \leftarrow b_a + b$$

This update should happen once for every example in every epoch, *irrespective of whether the weights were updated or not for that example*. In the end, the learning algorithm should return the averaged weights and the averaged bias.

(Technically, this strategy can be used with any of the variants we have seen here. For this homework, we only ask you to implement the averaged version of the original Perceptron. However, you are welcome to experiment with averaging the other variants.)

5. **Aggressive Perceptron with Margin, (For 6350 Students)** This algorithm is an extension of the margin Perceptron and performs an aggressive update as follows:

If $y(\mathbf{w}^T \mathbf{x}) \leq \mu$ then

$$(a) \quad \text{Update } \mathbf{w}_{new} \leftarrow \mathbf{w}_{old} + \eta y \mathbf{x}$$

Unlike the standard Perceptron algorithm, here the learning rate η is given by

$$\eta = \frac{\mu - y(\mathbf{w}^T \mathbf{x})}{\mathbf{x}^T \mathbf{x} + 1}$$

As with the margin Perceptron, there is an additional positive parameter μ .

Explanation of the update. We call this the aggressive update because the learning rate is derived from the following optimization problem:

When we see that $y(\mathbf{w}^T \mathbf{x}) \leq \mu$, we try to find new values of \mathbf{w} such that $y(\mathbf{w}^T \mathbf{x}) = \mu$ using

$$\begin{aligned} \min_{\mathbf{w}_{new}} \quad & \frac{1}{2} \|\mathbf{w}_{new} - \mathbf{w}_{old}\|^2 \\ \text{such that} \quad & y(\mathbf{w}^T \mathbf{x}) = \mu. \end{aligned}$$

That is, the goal is to find the smallest change in the weights so that the current example is on the right side of the weight vector.

By substituting (a) from above into this optimization problem, we will get a single variable optimization problem whose solution gives us the η defined above. You can think of this algorithm as trying to tune the weight vector so that the current example is correctly classified right after the update.

Implement this aggressive Perceptron algorithm.

Hyper-parameters: $\mu \in \{1, 0.1, 0.01\}$

4.3 Experiments

For all 5 settings above (4 for undergraduate students), you need to do the following things:

1. Run cross validation for **ten** epochs for each hyper-parameter combination to get the best hyper-parameter setting. Note that for cases when you are exploring combinations of hyper-parameters (such as the margin Perceptron), you need to try out all combinations.
2. Train the classifier for **20** epochs. At the end of each training epoch, you should measure the accuracy of the classifier on the development set. For the averaged Perceptron, use the average classifier to compute accuracy.
3. Use the classifier from the epoch where the development set accuracy is highest to evaluate on the test set.

4.4 What to report

1. [8 points] Briefly describe the design decisions that you have made in your implementation. (E.g, what programming language, how do you represent the vectors, etc.)
2. [2 points] *Majority baseline:* Consider a classifier that always predicts the most frequent label. What is its accuracy on test and development set?
3. [10 points per variant] For each variant above (5 for 6350 students, 4 for 5350 students), you need to report:
 - (a) The best hyper-parameters
 - (b) The cross-validation accuracy for the best hyperparameter
 - (c) The total number of updates the learning algorithm performs on the training set
 - (d) Development set accuracy
 - (e) Test set accuracy
 - (f) Plot a *learning curve* where the x axis is the epoch id and the y axis is the dev set accuracy using the classifier (or the averaged classifier, as appropriate) at the end of that epoch. Note that you should have selected the number of epochs using the learning curve (but no more than 20 epochs).

Experiment Submission Guidelines

1. The report should detail your experiments. For each step, explain in no more than a paragraph or so how your implementation works. Describe what you did. Comment on the design choices in your implementation. For your experiments, what algorithm parameters did you use? Try to analyze this and give your observations.
2. Your report should be in the form of a *pdf* file, \LaTeX is recommended.
3. *Your code should run on the CADE machines.* You should include a shell script, `run.sh`, that will execute your code in the CADE environment. Your code should produce similar output to what you include in your report.

You are responsible for ensuring that the grader can execute the code using only the included script. If you are using an esoteric programming language, you should make sure that its runtime is available on CADE.

4. Please do not hand in binary files! We will *not* grade binary submissions.
5. Please look up the late policy on the course website.